

Predicting Daily Cases Per Country

In this notebook data of a number of countries is used to train a model per each country to be used to predict the daily cases on a specific day.

In [0]:

```
import numpy as np
import pickle
import pandas as pd
from pandas import DataFrame
```

In [0]:

```
countries = {"Germany", "France", "Italy"}
```

Loading the data

In [0]:

```
def load_pickle(file_name):
    with open(file_name, 'rb') as f:
        return pickle.load(f)
```

Mounting Google Drive

We mount google drive to access the data stored there

In [181]:

```
from google.colab import drive
drive.mount('./gdrive')
```

Drive already mounted at ./gdrive; to attempt to forcibly remount, call drive.mount("./gdrive", force_remount=True).

In [0]:

```
drive_base_path = "./gdrive/My Drive"
data_path = "{} /COVID-19".format(drive_base_path)
```

Loading the Weather Data of the Countries and Joining Them

In [0]:

```
weather_data_base_path = "{} /weather-features".format(data_path)
```

In [0]:

```
wind_speed_dict = dict(filter(lambda x: x[0] in countries, load_pickle("{} /windspeedKmph_dict.pickle".format(
    weather_data_base_path)).items()))
tempreture_dict = dict(filter(lambda x: x[0] in countries, load_pickle("{} /tempC_dict.pickle".format(weather
    _data_base_path)).items()))
humidity_dict = dict(filter(lambda x: x[0] in countries, load_pickle("{} /humidity_dict.pickle".format(weathe
    r_data_base_path)).items()))
sun_hour_dict = dict(filter(lambda x: x[0] in countries, load_pickle("{} /sunHour_dict.pickle".format(weather
    _data_base_path)).items()))
```

In [0]:

```
germany_df = DataFrame()
germany_df['wind_speed'] = wind_speed_dict['Germany']
germany_df['tempreture'] = tempreture_dict['Germany']
germany_df['humidity'] = humidity_dict['Germany']
germany_df['sun_hour'] = sun_hour_dict['Germany']
```

In [0]:

```
italy_df = DataFrame()
italy_df['wind_speed'] = wind_speed_dict['Italy']
italy_df['tempreture'] = tempreture_dict['Italy']
italy_df['humidity'] = humidity_dict['Italy']
italy_df['sun_hour'] = sun_hour_dict['Italy']
```

In [0]:

```
france_df = DataFrame()
france_df['wind_speed'] = wind_speed_dict['France']
france_df['tempreture'] = tempreture_dict['France']
france_df['humidity'] = humidity_dict['France']
france_df['sun_hour'] = sun_hour_dict['France']
```

Loading "Our World in Data" Dataset and Merging it With the Weather Data

In [0]:

```
our_world_data_base_path = "{}our-world-in-data".format(data_path)
```

In [0]:

```
def merge_dataframes_filter_with_date(main_df, to_be_merged_df, column_names, start_date, end_date):
    to_be_merged_df['date'] = pd.to_datetime(to_be_merged_df['date'])
    to_be_merged_df = to_be_merged_df[(to_be_merged_df['date'] >= start_date) & (to_be_merged_df['date'] <= end_date)]
    to_be_merged_df.index = [x for x in range(main_df.shape[0])]
    main_df['date'] = to_be_merged_df['date']
    for name in column_names:
        main_df[name] = to_be_merged_df[name]
```

In [0]:

```
new_cases_dict = dict(filter(lambda x: x[0] in countries, load_pickle("{}new_cases_dict.pickle".format(our_world_data_base_path)).items()))
```

In [0]:

```
new_deaths_dict = dict(filter(lambda x: x[0] in countries, load_pickle("{}new_deaths_dict.pickle".format(our_world_data_base_path)).items()))
```

We have to restrict the days of the data to be the same as the days of weather data to be able to join them

In [0]:

```
start_date = '2020-01-22'
end_date = '2020-03-21'
```

In [0]:

```
germany_new_cases_df = DataFrame(new_cases_dict['Germany'], columns=['date', 'new_cases'])
merge_dataframes_filter_with_date(germany_df, germany_new_cases_df, ['new_cases'], start_date, end_date)
germany_new_deaths_df = DataFrame(new_deaths_dict['Germany'], columns=['date', 'new_deaths'])
merge_dataframes_filter_with_date(germany_df, germany_new_deaths_df, ['new_deaths'], start_date, end_date)
```

In [0]:

```
italy_new_cases_df = DataFrame(new_cases_dict['Italy'], columns=['date', 'new_cases'])
merge_dataframes_filter_with_date(italy_df, italy_new_cases_df, ['new_cases'], start_date, end_date)
italy_new_deaths_df = DataFrame(new_deaths_dict['Italy'], columns=['date', 'new_deaths'])
merge_dataframes_filter_with_date(italy_df, italy_new_deaths_df, ['new_deaths'], start_date, end_date)
```

In [0]:

```
france_new_cases_df = DataFrame(new_cases_dict['France'], columns=['date', 'new_cases'])
merge_dataframes_filter_with_date(france_df, france_new_cases_df, ['new_cases'], start_date, end_date)
france_new_deaths_df = DataFrame(new_deaths_dict['France'], columns=['date', 'new_deaths'])
merge_dataframes_filter_with_date(france_df, france_new_deaths_df, ['new_deaths'], start_date, end_date)
```

Defining the Models

In this section we are going to try 2 approaches.

The first one is to consider the data as a time series and to use LSTM in the model and train it to predict the cases of a day based on a previous day. The time series forecasting problem can be trained to have different types of inputs and outputs. One approach is to input a day and output a day. A second approach is to have an input sequence of days and output a day. Another approach is to have a input sequence of days and output another sequence of days. We are going to use the first approach with the input day as one day before the output day. The number of days the input is before the output day is a hyper parameter that can be changed to get the best results. This [paper](https://www.researchgate.net/publication/341089678_Neural_Network_Model_for_Prediction_of_Covid-19_Confirmed_Cases_and_Fatalities)


(https://www.researchgate.net/publication/341089678_Neural_Network_Model_for_Prediction_of_Covid-19_Confirmed_Cases_and_Fatalities) is used as a reference for the parameters used in the LSTM model.

The second approach we are going to use is a fully connected neural network with one hidden layer. We use root mean square error as the error function to evaluate the models.

The models are tested using Germany, France and Italy's data.

LSTM

Long short-term memory (LSTM) is an artificial recurrent neural network (RNN) architecture. A common LSTM unit is composed of a cell, an input gate, an output gate and a forget gate. The cell remembers values over arbitrary time intervals and the three gates regulate the flow of information into and out of the cell. LSTM networks are well-suited to classifying, processing and making predictions based on time series data, since there can be lags of unknown duration between important events in a time series. Intuitively, the cell is responsible for keeping track of the dependencies between the elements in the input sequence. The input gate controls the extent to which a new value flows into the cell, the forget gate controls the extent to which a value remains in the cell and the output gate controls the extent to which the value in the cell is used to compute the output activation of the LSTM unit. The activation function of the LSTM gates is often the logistic sigmoid function.



In [0]:

```
import math
import keras
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import LSTM
from keras.layers import Dropout
from keras.layers import *
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
from keras.callbacks import EarlyStopping
```

Germany's Models

In [0]:

```
germany_df.drop(columns='date', inplace=True)
```

The LSTM Model

Preparing the train and Test Sets

First thing to do is to add a new column which is a shift by one day of the new cases model to be used as the output for the LSTM model for the training phase

In [0]:

```
germany_df['output_new_cases'] = germany_df['new_cases'].shift(-1)
germany_df.drop(axis=0, index=[germany_df.shape[0]-1], inplace=True)
```

We split the data into a training, test and validation sets 70%, 10%, 20% splits respectively, after normalizing the data using MinMaxScaler in order to remove the dominance of large valued features. The sets have to be reshaped to [examples, timesteps, features] to be used as inputs to the LSTM model. In our case the timesteps used is 1 as only 1 day is considered as input not a sequence of days

In [0]:

```
germany_dataset = germany_df.values
germany_dataset = germany_dataset.astype('float32')
scaler = MinMaxScaler(feature_range=(0, 1))
germany_dataset = scaler.fit_transform(germany_dataset)
train_size = int(len(germany_dataset) * 0.70)
val_size = int(len(germany_dataset) * 0.10)
test_size = int(len(germany_dataset) * 0.20)
train, val, test = germany_dataset[0:train_size,:], germany_dataset[train_size:val_size + train_size,:], germany_dataset[val_size + train_size:len(germany_dataset),:]
train_X, train_Y = train[:, :-1], train[:, -1]
test_X, test_Y = test[:, :-1], test[:, -1]
val_X, val_Y = val[:, :-1], val[:, -1]
train_X = np.reshape(train_X, (train_X.shape[0], 1, train_X.shape[1]))
test_X = np.reshape(test_X, (test_X.shape[0], 1, test_X.shape[1]))
val_X = np.reshape(val_X, (val_X.shape[0], 1, val_X.shape[1]))
```

Training the LSTM Model

In [244]:

```
lstm_model = Sequential()
lstm_model.add(LSTM(64, input_shape=(train_X.shape[1], train_X.shape[2])))
lstm_model.add(Dropout(0.5))
lstm_model.add(Dense(1))
lstm_model.compile(loss='mean_squared_error', optimizer='adam')

history = lstm_model.fit(train_X, train_Y, epochs=20, batch_size=70, verbose=1, shuffle=False,
                        validation_data=(val_X, val_Y), callbacks=[EarlyStopping(monitor='val_loss', patience=10)])
lstm_model.summary()
```

Train on 41 samples, validate on 5 samples

```
Epoch 1/20
41/41 [=====] - 0s 8ms/step - loss: 0.0018 - val_loss: 7.6273e-04
Epoch 2/20
41/41 [=====] - 0s 147us/step - loss: 0.0015 - val_loss: 0.0013
Epoch 3/20
41/41 [=====] - 0s 125us/step - loss: 0.0018 - val_loss: 0.0018
Epoch 4/20
41/41 [=====] - 0s 91us/step - loss: 0.0019 - val_loss: 0.0022
Epoch 5/20
41/41 [=====] - 0s 97us/step - loss: 0.0025 - val_loss: 0.0022
Epoch 6/20
41/41 [=====] - 0s 94us/step - loss: 0.0014 - val_loss: 0.0022
Epoch 7/20
41/41 [=====] - 0s 104us/step - loss: 0.0015 - val_loss: 0.0021
Epoch 8/20
41/41 [=====] - 0s 90us/step - loss: 0.0012 - val_loss: 0.0019
Epoch 9/20
41/41 [=====] - 0s 97us/step - loss: 9.8731e-04 - val_loss: 0.0016
Epoch 10/20
41/41 [=====] - 0s 89us/step - loss: 0.0010 - val_loss: 0.0014
Epoch 11/20
41/41 [=====] - 0s 97us/step - loss: 0.0013 - val_loss: 0.0012
Model: "sequential_24"
```

Layer (type)	Output Shape	Param #
lstm_9 (LSTM)	(None, 64)	18432
dropout_19 (Dropout)	(None, 64)	0
dense_31 (Dense)	(None, 1)	65
Total params: 18,497		
Trainable params: 18,497		
Non-trainable params: 0		

Predicting Using the LSTM Model

In [245]:

```
# Make a prediction
yhat = lstm_model.predict(test_X)
test_X = test_X.reshape((test_X.shape[0], test_X.shape[2]))
test_Y = test_Y.reshape((len(test_Y), 1))

# Calculate RMSE
rmse = np.sqrt(mean_squared_error(test_Y, yhat))
print('Test RMSE: %.3f' % rmse)
```

Test RMSE: 0.309

The Dense Model

Preparing the Train and Test Sets

The data is split and normalized as done before the LSTM model but not reshaped as the input doesn't contain the timesteps dimension.

In [0]:

```
dense_model_germany_df = germany_df
#Add a counter as a feature to indicate the day from the start of the pandemic
dense_model_germany_df['days_from_pandemic_start'] = [x for x in range(dense_model_germany_df.shape[0])]
g_new_cases_df = dense_model_germany_df['new_cases']
dense_model_germany_df.drop(columns='new_cases', inplace=True)
dense_model_germany_df['new_cases'] = g_new_cases_df
germany_dense_model_dataset = dense_model_germany_df.values
germany_dense_model_dataset = germany_dense_model_dataset.astype('float32')
dense_model_scaler = MinMaxScaler(feature_range=(0, 1))
germany_dense_model_dataset = dense_model_scaler.fit_transform(germany_dense_model_dataset)
train_size = int(len(germany_dataset) * 0.70)
val_size = int(len(germany_dataset) * 0.10)
test_size = int(len(germany_dataset) * 0.20)
train, val, test = germany_dense_model_dataset[0:train_size,:], germany_dense_model_dataset[train_size:val_size + train_size,:], germany_dense_model_dataset[val_size + train_size:len(germany_dataset),:]
train_X, train_Y = train[:, :-1], train[:, -1]
test_X, test_Y = test[:, :-1], test[:, -1]
val_X, val_Y = val[:, :-1], val[:, -1]
```

Defining the Dense Model

In [250]:

```
model = Sequential()
model.add(Dense(64, input_dim=train_X.shape[1], activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(1, activation='linear'))
model.compile(loss='mean_squared_error', optimizer='adam')

history = model.fit(train_X, train_Y, epochs=20, batch_size=70, verbose=1, shuffle=False,
                    validation_data=(val_X, val_Y), callbacks=[EarlyStopping(monitor='val_loss', patience=10)])

model.summary()
```

Train on 41 samples, validate on 5 samples

Epoch 1/20
41/41 [=====] - 0s 2ms/step - loss: 0.0706 - val_loss: 0.0497

Epoch 2/20
41/41 [=====] - 0s 124us/step - loss: 0.0712 - val_loss: 0.0421

Epoch 3/20
41/41 [=====] - 0s 97us/step - loss: 0.0471 - val_loss: 0.0349

Epoch 4/20
41/41 [=====] - 0s 104us/step - loss: 0.0694 - val_loss: 0.0278

Epoch 5/20
41/41 [=====] - 0s 97us/step - loss: 0.0540 - val_loss: 0.0220

Epoch 6/20
41/41 [=====] - 0s 101us/step - loss: 0.0663 - val_loss: 0.0168

Epoch 7/20
41/41 [=====] - 0s 102us/step - loss: 0.0579 - val_loss: 0.0126

Epoch 8/20
41/41 [=====] - 0s 56us/step - loss: 0.0401 - val_loss: 0.0095

Epoch 9/20
41/41 [=====] - 0s 59us/step - loss: 0.0221 - val_loss: 0.0073

Epoch 10/20
41/41 [=====] - 0s 108us/step - loss: 0.0332 - val_loss: 0.0059

Epoch 11/20
41/41 [=====] - 0s 58us/step - loss: 0.0424 - val_loss: 0.0052

Epoch 12/20
41/41 [=====] - 0s 61us/step - loss: 0.0367 - val_loss: 0.0051

Epoch 13/20
41/41 [=====] - 0s 54us/step - loss: 0.0199 - val_loss: 0.0056

Epoch 14/20
41/41 [=====] - 0s 69us/step - loss: 0.0332 - val_loss: 0.0063

Epoch 15/20
41/41 [=====] - 0s 56us/step - loss: 0.0250 - val_loss: 0.0073

Epoch 16/20
41/41 [=====] - 0s 58us/step - loss: 0.0224 - val_loss: 0.0084

Epoch 17/20
41/41 [=====] - 0s 57us/step - loss: 0.0219 - val_loss: 0.0096

Epoch 18/20
41/41 [=====] - 0s 60us/step - loss: 0.0182 - val_loss: 0.0110

Epoch 19/20
41/41 [=====] - 0s 75us/step - loss: 0.0209 - val_loss: 0.0122

Epoch 20/20
41/41 [=====] - 0s 60us/step - loss: 0.0197 - val_loss: 0.0133

Model: "sequential_25"

Layer (type)	Output Shape	Param #
dense_32 (Dense)	(None, 64)	512
dropout_20 (Dropout)	(None, 64)	0
dense_33 (Dense)	(None, 1)	65
Total params: 577		
Trainable params: 577		
Non-trainable params: 0		

Predicting Using the Dense Model

In [251]:

```
# make a prediction
yhat = model.predict(test_X)
test_Y = test_Y.reshape((len(test_Y), 1))

# calculate RMSE
rmse = np.sqrt(mean_squared_error(test_Y, yhat))
print('Test RMSE: %.3f' % rmse)
```

Test RMSE: 0.231

The same steps are repeated for the France and the Italy's datasets.

France's Models

In [0]:

```
france_df.drop(columns='date', inplace=True)
```

The LSTM Model

Preparing the train and Test Sets

In [0]:

```
france_df['output_new_cases'] =france_df['new_cases'].shift(-1)
france_df.drop(axis=0,index=[france_df.shape[0]-1], inplace=True)
```

In [0]:

```
france_dataset = france_df.values
france_dataset = france_dataset.astype('float32')
scaler = MinMaxScaler(feature_range=(0, 1))
france_dataset = scaler.fit_transform(france_dataset)
train_size = int(len(france_dataset) * 0.70)
val_size = int(len(france_dataset) * 0.10)
test_size = int(len(france_dataset) * 0.20)
train, val ,test = france_dataset[0:train_size,:], france_dataset[train_size:val_size + train_size,:], franc
e_dataset[val_size + train_size:len(france_dataset),:]
train_X, train_Y = train[:,-1], train[:, -1]
test_X, test_Y = test[:, :-1], test[:, -1]
val_X, val_Y = val[:, :-1], val[:, -1]
train_X = np.reshape(train_X, (train_X.shape[0], 1, train_X.shape[1]))
test_X = np.reshape(test_X, (test_X.shape[0], 1, test_X.shape[1]))
val_X = np.reshape(val_X, (val_X.shape[0], 1, val_X.shape[1]))
```

Defining the LSTM Model

In [271]:

```
lstm_model = Sequential()
lstm_model.add(LSTM(64, input_shape=(train_X.shape[1], train_X.shape[2])))
lstm_model.add(Dropout(0.5))
lstm_model.add(Dense(1))
lstm_model.compile(loss='mean_squared_error', optimizer='adam')

history = lstm_model.fit(train_X, train_Y, epochs=20, batch_size=70, verbose=1, shuffle=False,
                        validation_data=(val_X, val_Y), callbacks=[EarlyStopping(monitor='val_loss', patien
ce=10)])
lstm_model.summary()
```


Train on 41 samples, validate on 5 samples

Epoch 1/20
41/41 [=====] - 0s 8ms/step - loss: 0.0015 - val_loss: 0.0048

Epoch 2/20
41/41 [=====] - 0s 147us/step - loss: 0.0011 - val_loss: 0.0043

Epoch 3/20
41/41 [=====] - 0s 123us/step - loss: 8.8675e-04 - val_loss: 0.0038

Epoch 4/20
41/41 [=====] - 0s 122us/step - loss: 6.7794e-04 - val_loss: 0.0034

Epoch 5/20
41/41 [=====] - 0s 117us/step - loss: 9.4882e-04 - val_loss: 0.0033

Epoch 6/20
41/41 [=====] - 0s 120us/step - loss: 6.8105e-04 - val_loss: 0.0032

Epoch 7/20
41/41 [=====] - 0s 119us/step - loss: 8.0140e-04 - val_loss: 0.0032

Epoch 8/20
41/41 [=====] - 0s 125us/step - loss: 6.9709e-04 - val_loss: 0.0033

Epoch 9/20
41/41 [=====] - 0s 119us/step - loss: 7.9482e-04 - val_loss: 0.0034

Epoch 10/20
41/41 [=====] - 0s 112us/step - loss: 6.1630e-04 - val_loss: 0.0036

Epoch 11/20
41/41 [=====] - 0s 121us/step - loss: 7.0715e-04 - val_loss: 0.0038

Epoch 12/20
41/41 [=====] - 0s 115us/step - loss: 5.5393e-04 - val_loss: 0.0041

Epoch 13/20
41/41 [=====] - 0s 114us/step - loss: 4.4567e-04 - val_loss: 0.0043

Epoch 14/20
41/41 [=====] - 0s 115us/step - loss: 7.5241e-04 - val_loss: 0.0044

Epoch 15/20
41/41 [=====] - 0s 114us/step - loss: 5.2984e-04 - val_loss: 0.0045

Epoch 16/20
41/41 [=====] - 0s 106us/step - loss: 5.0908e-04 - val_loss: 0.0045

Epoch 17/20
41/41 [=====] - 0s 93us/step - loss: 5.5986e-04 - val_loss: 0.0043

Model: "sequential_28"

Layer (type)	Output Shape	Param #
lstm_12 (LSTM)	(None, 64)	18176
dropout_23 (Dropout)	(None, 64)	0
dense_36 (Dense)	(None, 1)	65
Total params: 18,241		
Trainable params: 18,241		
Non-trainable params: 0		

Predicting Using the LSTM Model

In [272]:

```
# make a prediction
yhat = lstm_model.predict(test_X)
test_X = test_X.reshape((test_X.shape[0], test_X.shape[2]))
test_Y = test_Y.reshape((len(test_Y), 1))

# calculate RMSE
rmse = np.sqrt(mean_squared_error(test_Y, yhat))
print('Test RMSE: %.3f' % rmse)
```

Test RMSE: 0.572

The Dense Model

Preparing the Train and Test Sets

In [0]:

```
dense_model_france_df = france_df
#Add a counter as a feature to indicate the day from the start of the pandemic
dense_model_france_df['days_from_pandemic_start'] = [x for x in range(dense_model_france_df.shape[0])]
g_new_cases_df = dense_model_france_df['new_cases']
dense_model_france_df.drop(columns='new_cases', inplace=True)
dense_model_france_df['new_cases'] = g_new_cases_df
france_dense_model_dataset = dense_model_france_df.values
france_dense_model_dataset = france_dense_model_dataset.astype('float32')
dense_model_scaler = MinMaxScaler(feature_range=(0, 1))
france_dense_model_dataset = dense_model_scaler.fit_transform(france_dense_model_dataset)
train_size = int(len(france_dataset) * 0.70)
val_size = int(len(france_dataset) * 0.10)
test_size = int(len(france_dataset) * 0.20)
train, val, test = france_dense_model_dataset[0:train_size,:], france_dense_model_dataset[train_size:val_size + train_size:], france_dense_model_dataset[val_size + train_size:len(france_dataset),:]
train_X, train_Y = train[:, :-1], train[:, -1]
test_X, test_Y = test[:, :-1], test[:, -1]
val_X, val_Y = val[:, :-1], val[:, -1]
```

Defining the Dense Model

In [275]:

```
model = Sequential()
model.add(Dense(64, input_dim=train_X.shape[1], activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(1, activation='linear'))
model.compile(loss='mean_squared_error', optimizer='adam')

history = model.fit(train_X, train_Y, epochs=20, batch_size=70, verbose=1, shuffle=False,
                    validation_data=(val_X, val_Y), callbacks=[EarlyStopping(monitor='val_loss', patience=10)])

model.summary()
```

Train on 41 samples, validate on 5 samples

Epoch 1/20
41/41 [=====] - 0s 2ms/step - loss: 0.0481 - val_loss: 0.0164

Epoch 2/20
41/41 [=====] - 0s 86us/step - loss: 0.0414 - val_loss: 0.0128

Epoch 3/20
41/41 [=====] - 0s 65us/step - loss: 0.0371 - val_loss: 0.0102

Epoch 4/20
41/41 [=====] - 0s 70us/step - loss: 0.0397 - val_loss: 0.0083

Epoch 5/20
41/41 [=====] - 0s 53us/step - loss: 0.0212 - val_loss: 0.0070

Epoch 6/20
41/41 [=====] - 0s 59us/step - loss: 0.0224 - val_loss: 0.0060

Epoch 7/20
41/41 [=====] - 0s 60us/step - loss: 0.0227 - val_loss: 0.0055

Epoch 8/20
41/41 [=====] - 0s 57us/step - loss: 0.0214 - val_loss: 0.0053

Epoch 9/20
41/41 [=====] - 0s 60us/step - loss: 0.0262 - val_loss: 0.0052

Epoch 10/20
41/41 [=====] - 0s 56us/step - loss: 0.0264 - val_loss: 0.0053

Epoch 11/20
41/41 [=====] - 0s 60us/step - loss: 0.0152 - val_loss: 0.0055

Epoch 12/20
41/41 [=====] - 0s 63us/step - loss: 0.0178 - val_loss: 0.0057

Epoch 13/20
41/41 [=====] - 0s 73us/step - loss: 0.0191 - val_loss: 0.0059

Epoch 14/20
41/41 [=====] - 0s 65us/step - loss: 0.0222 - val_loss: 0.0059

Epoch 15/20
41/41 [=====] - 0s 72us/step - loss: 0.0264 - val_loss: 0.0059

Epoch 16/20
41/41 [=====] - 0s 69us/step - loss: 0.0199 - val_loss: 0.0057

Epoch 17/20
41/41 [=====] - 0s 68us/step - loss: 0.0160 - val_loss: 0.0054

Epoch 18/20
41/41 [=====] - 0s 72us/step - loss: 0.0245 - val_loss: 0.0051

Epoch 19/20
41/41 [=====] - 0s 76us/step - loss: 0.0208 - val_loss: 0.0047

Epoch 20/20
41/41 [=====] - 0s 81us/step - loss: 0.0115 - val_loss: 0.0043

Model: "sequential_30"

Layer (type)	Output Shape	Param #
dense_39 (Dense)	(None, 64)	512
dropout_25 (Dropout)	(None, 64)	0
dense_40 (Dense)	(None, 1)	65

Total params: 577
Trainable params: 577
Non-trainable params: 0

Predicting Using the Dense Model

In [276]:

```
# make a prediction
yhat = model.predict(test_X)
test_Y = test_Y.reshape((len(test_Y), 1))

# calculate RMSE
rmse = np.sqrt(mean_squared_error(test_Y, yhat))
print('Test RMSE: %.3f' % rmse)
```

Test RMSE: 0.316

In [0]:

Italy's Models

In [0]:

```
italy_df.drop(columns='date', inplace=True)
```

The LSTM Model

Preparing the train and Test Sets

In [0]:

```
italy_df['output_new_cases'] = italy_df['new_cases'].shift(-1)
italy_df.drop(axis=0, index=[italy_df.shape[0]-1], inplace=True)
```

In [0]:

```
italy_dataset = france_df.values
italy_dataset = italy_dataset.astype('float32')
scaler = MinMaxScaler(feature_range=(0, 1))
italy_dataset = scaler.fit_transform(italy_dataset)
train_size = int(len(italy_dataset) * 0.70)
val_size = int(len(italy_dataset) * 0.10)
test_size = int(len(italy_dataset) * 0.20)
train, val, test = italy_dataset[0:train_size,:], italy_dataset[train_size:val_size + train_size,:], italy_dataset[val_size + train_size:len(italy_dataset),:]
train_X, train_Y = train[:, :-1], train[:, -1]
test_X, test_Y = test[:, :-1], test[:, -1]
val_X, val_Y = val[:, :-1], val[:, -1]
train_X = np.reshape(train_X, (train_X.shape[0], 1, train_X.shape[1]))
test_X = np.reshape(test_X, (test_X.shape[0], 1, test_X.shape[1]))
val_X = np.reshape(val_X, (val_X.shape[0], 1, val_X.shape[1]))
```

Defining the LSTM Model

In [280]:

```
lstm_model = Sequential()
lstm_model.add(LSTM(64, input_shape=(train_X.shape[1], train_X.shape[2])))
lstm_model.add(Dropout(0.5))
lstm_model.add(Dense(1))
lstm_model.compile(loss='mean_squared_error', optimizer='adam')

history = lstm_model.fit(train_X, train_Y, epochs=20, batch_size=70, verbose=1, shuffle=False,
                        validation_data=(val_X, val_Y), callbacks=[EarlyStopping(monitor='val_loss', patience=10)])
lstm_model.summary()
```

Train on 41 samples, validate on 5 samples

Epoch 1/20
41/41 [=====] - 0s 8ms/step - loss: 0.0015 - val_loss: 0.0032

Epoch 2/20
41/41 [=====] - 0s 149us/step - loss: 0.0024 - val_loss: 0.0032

Epoch 3/20
41/41 [=====] - 0s 266us/step - loss: 0.0024 - val_loss: 0.0033

Epoch 4/20
41/41 [=====] - 0s 145us/step - loss: 0.0011 - val_loss: 0.0033

Epoch 5/20
41/41 [=====] - 0s 139us/step - loss: 0.0012 - val_loss: 0.0032

Epoch 6/20
41/41 [=====] - 0s 109us/step - loss: 0.0011 - val_loss: 0.0031

Epoch 7/20
41/41 [=====] - 0s 105us/step - loss: 0.0012 - val_loss: 0.0032

Epoch 8/20
41/41 [=====] - 0s 110us/step - loss: 0.0012 - val_loss: 0.0031

Epoch 9/20
41/41 [=====] - 0s 121us/step - loss: 9.4214e-04 - val_loss: 0.0030

Epoch 10/20
41/41 [=====] - 0s 117us/step - loss: 0.0010 - val_loss: 0.0030

Epoch 11/20
41/41 [=====] - 0s 116us/step - loss: 0.0011 - val_loss: 0.0028

Epoch 12/20
41/41 [=====] - 0s 108us/step - loss: 8.6113e-04 - val_loss: 0.0027

Epoch 13/20
41/41 [=====] - 0s 110us/step - loss: 8.3132e-04 - val_loss: 0.0026

Epoch 14/20
41/41 [=====] - 0s 109us/step - loss: 0.0011 - val_loss: 0.0027

Epoch 15/20
41/41 [=====] - 0s 88us/step - loss: 7.8429e-04 - val_loss: 0.0028

Epoch 16/20
41/41 [=====] - 0s 82us/step - loss: 6.3349e-04 - val_loss: 0.0030

Epoch 17/20
41/41 [=====] - 0s 112us/step - loss: 8.1473e-04 - val_loss: 0.0031

Epoch 18/20
41/41 [=====] - 0s 98us/step - loss: 0.0010 - val_loss: 0.0032

Epoch 19/20
41/41 [=====] - 0s 97us/step - loss: 7.0902e-04 - val_loss: 0.0032

Epoch 20/20
41/41 [=====] - 0s 124us/step - loss: 6.8554e-04 - val_loss: 0.0032

Model: "sequential_31"

Layer (type)	Output Shape	Param #
lstm_13 (LSTM)	(None, 64)	18432
dropout_26 (Dropout)	(None, 64)	0
dense_41 (Dense)	(None, 1)	65
Total params: 18,497		
Trainable params: 18,497		
Non-trainable params: 0		

Predicting Using the LSTM Model

In [281]:

```
# make a prediction
yhat = lstm_model.predict(test_X)
test_X = test_X.reshape((test_X.shape[0], test_X.shape[2]))
test_Y = test_Y.reshape((len(test_Y), 1))

# calculate RMSE
rmse = np.sqrt(mean_squared_error(test_Y, yhat))
print('Test RMSE: %.3f' % rmse)
```

Test RMSE: 0.493

The Dense Model

Preparing the Train and Test Sets

In [0]:

```
dense_model_italy_df = italy_df
#Add a counter as a feature to indicate the day from the start of the pandemic
dense_model_italy_df['days_from_pandemic_start'] = [x for x in range(dense_model_italy_df.shape[0])]
g_new_cases_df = dense_model_italy_df['new_cases']
dense_model_italy_df.drop(columns='new_cases', inplace=True)
dense_model_italy_df['new_cases'] = g_new_cases_df
italy_dense_model_dataset = dense_model_italy_df.values
italy_dense_model_dataset = italy_dense_model_dataset.astype('float32')
dense_model_scaler = MinMaxScaler(feature_range=(0, 1))
italy_dense_model_dataset = dense_model_scaler.fit_transform(italy_dense_model_dataset)
train_size = int(len(italy_dataset) * 0.70)
val_size = int(len(italy_dataset) * 0.10)
test_size = int(len(italy_dataset) * 0.20)
train, val, test = italy_dense_model_dataset[0:train_size,:], italy_dense_model_dataset[train_size:val_size
+ train_size,:], italy_dense_model_dataset[val_size + train_size:len(italy_dataset),:]
train_X, train_Y = train[:, :-1], train[:, -1]
test_X, test_Y = test[:, :-1], test[:, -1]
val_X, val_Y = val[:, :-1], val[:, -1]
```

Defining the Dense Model

In [290]:

```
model = Sequential()
model.add(Dense(64, input_dim=train_X.shape[1], activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(1, activation='linear'))
model.compile(loss='mean_squared_error', optimizer='adam')

history = model.fit(train_X, train_Y, epochs=20, batch_size=70, verbose=1, shuffle=False,
                    validation_data=(val_X, val_Y), callbacks=[EarlyStopping(monitor='val_loss', patience=10
)])

model.summary()
```

Train on 41 samples, validate on 5 samples

Epoch 1/20
41/41 [=====] - 0s 2ms/step - loss: 0.0468 - val_loss: 0.1060

Epoch 2/20
41/41 [=====] - 0s 69us/step - loss: 0.0552 - val_loss: 0.0952

Epoch 3/20
41/41 [=====] - 0s 99us/step - loss: 0.0517 - val_loss: 0.0851

Epoch 4/20
41/41 [=====] - 0s 80us/step - loss: 0.0436 - val_loss: 0.0764

Epoch 5/20
41/41 [=====] - 0s 65us/step - loss: 0.0197 - val_loss: 0.0692

Epoch 6/20
41/41 [=====] - 0s 66us/step - loss: 0.0359 - val_loss: 0.0623

Epoch 7/20
41/41 [=====] - 0s 72us/step - loss: 0.0334 - val_loss: 0.0556

Epoch 8/20
41/41 [=====] - 0s 64us/step - loss: 0.0275 - val_loss: 0.0496

Epoch 9/20
41/41 [=====] - 0s 60us/step - loss: 0.0300 - val_loss: 0.0442

Epoch 10/20
41/41 [=====] - 0s 64us/step - loss: 0.0321 - val_loss: 0.0394

Epoch 11/20
41/41 [=====] - 0s 71us/step - loss: 0.0253 - val_loss: 0.0355

Epoch 12/20
41/41 [=====] - 0s 66us/step - loss: 0.0258 - val_loss: 0.0323

Epoch 13/20
41/41 [=====] - 0s 68us/step - loss: 0.0230 - val_loss: 0.0295

Epoch 14/20
41/41 [=====] - 0s 69us/step - loss: 0.0269 - val_loss: 0.0272

Epoch 15/20
41/41 [=====] - 0s 70us/step - loss: 0.0354 - val_loss: 0.0253

Epoch 16/20
41/41 [=====] - 0s 65us/step - loss: 0.0158 - val_loss: 0.0237

Epoch 17/20
41/41 [=====] - 0s 67us/step - loss: 0.0188 - val_loss: 0.0224

Epoch 18/20
41/41 [=====] - 0s 68us/step - loss: 0.0283 - val_loss: 0.0216

Epoch 19/20
41/41 [=====] - 0s 68us/step - loss: 0.0199 - val_loss: 0.0210

Epoch 20/20
41/41 [=====] - 0s 68us/step - loss: 0.0146 - val_loss: 0.0206

Model: "sequential_34"

Layer (type)	Output Shape	Param #
dense_46 (Dense)	(None, 64)	512
dropout_29 (Dropout)	(None, 64)	0
dense_47 (Dense)	(None, 1)	65
Total params: 577		
Trainable params: 577		
Non-trainable params: 0		

Predicting Using the Dense Model

In [291]:

```
# make a prediction
yhat = model.predict(test_X)
test_Y = test_Y.reshape((len(test_Y), 1))

# calculate RMSE
rmse = np.sqrt(mean_squared_error(test_Y, yhat))
print('Test RMSE: %.3f' % rmse)
```

Test RMSE: 0.471

Conclusion

The results are close when using the LSTM model or the dense model with the dense model having slightly better results. The amount of the data was a limitation with only 60 days used as the weather data is the resitricting factor. Results would have been better and a better comparison would have been made if more data samples are used in training.

References

- <https://colah.github.io/posts/2015-08-Understanding-LSTMs/> (<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>)
- <https://towardsdatascience.com/time-series-analysis-visualization-forecasting-with-lstm-77a905180eba>
(<https://towardsdatascience.com/time-series-analysis-visualization-forecasting-with-lstm-77a905180eba>)
- <https://machinelearningmastery.com/multivariate-time-series-forecasting-lstms-keras/>
(<https://machinelearningmastery.com/multivariate-time-series-forecasting-lstms-keras/>)
- <https://machinelearningmastery.com/how-to-choose-loss-functions-when-training-deep-learning-neural-networks/>
(<https://machinelearningmastery.com/how-to-choose-loss-functions-when-training-deep-learning-neural-networks/>)
- <https://towardsdatascience.com/lstm-for-time-series-prediction-de8aeb26f2ca> (<https://towardsdatascience.com/lstm-for-time-series-prediction-de8aeb26f2ca>)
- <https://machinelearningmastery.com/convert-time-series-supervised-learning-problem-python/>
(<https://machinelearningmastery.com/convert-time-series-supervised-learning-problem-python/>)
- <https://machinelearningmastery.com/how-to-develop-lstm-models-for-time-series-forecasting/>
(<https://machinelearningmastery.com/how-to-develop-lstm-models-for-time-series-forecasting/>)